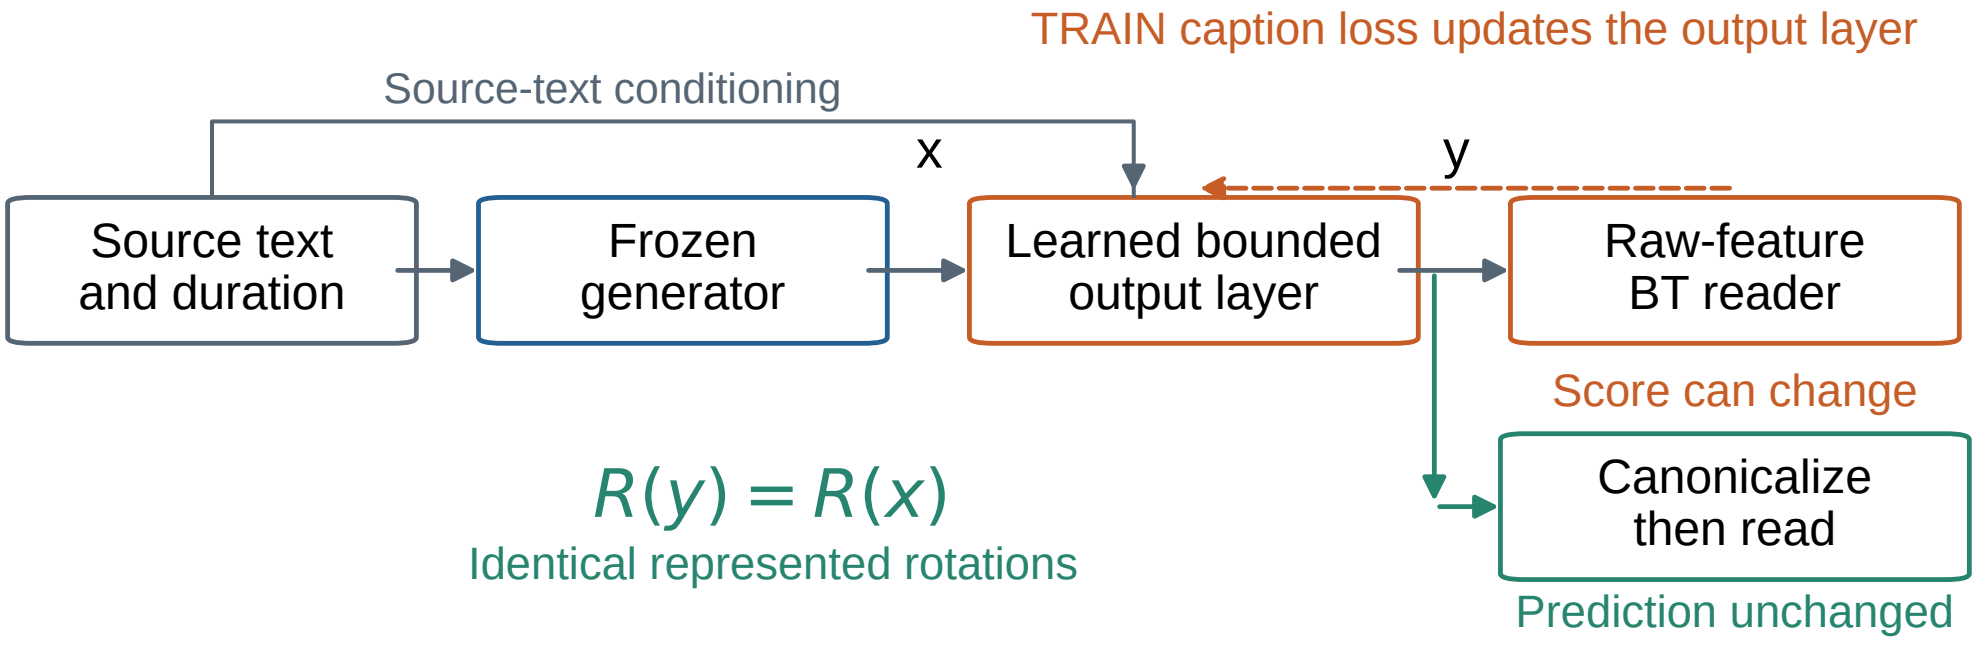
TRAIN caption loss updates the output layer
Source-text conditioning
x
y
Source text and duration
Frozen generator
Learned bounded output layer
Raw-feature BT reader
Score can change
Canonicalize then read
$R(y) = R(x)$
Identical represented rotations
Prediction unchanged